

Written Assessment Answers

1. Agent Isolation

A system prompt can guide an agent, but it should not be responsible for security. Each agent should receive only the data and approved functions it needs for its department. The model should never have direct database access or be allowed to run arbitrary SQL. In production, I would also enforce the separation with server side authorization, Supabase Row Level Security, and restricted database functions. Every request should check the user, company, and department before returning any data. If the permissions are unclear, access should be denied. In my implementation, the Finance and HR agents receive separate data and types, and the tests verify that they cannot access each other's fields.

2. AI Agent Meeting

I would use a central coordinator to manage the meeting instead of letting the agents freely talk to each other. Each agent would first retrieve verified data from its own department and return a structured summary. Every important number should include its source and reporting period. The coordinator would combine the summaries, but it would not be allowed to introduce new numbers. If two agents provide conflicting information, the data should be checked again rather than guessed. Missing information should also be clearly stated. I used the same approach in my project: both agents work independently, and the coordinator receives only their verified analyses.

3. Scale and Rate Limits

15 employees sending 2 to 3 requests per minute would create around 30 - 45 requests per minute. I would send all API requests through one shared client that controls concurrency. Interactive requests would have priority, while background work could wait in a queue. If a request fails because of rate limits or a temporary error, I would retry it after a short delay. I would also cache repeated requests when the data has not changed, and use smaller models for simple tasks while reserving stronger models for more complex analysis. In my implementation, the LLM client already includes retries, model selection, and a token budget..